Restore the missing history, preserve the resident pages

Conceptual layout: 16 tokens per GPU page; 16 pages per CPU chunk (256 tokens)

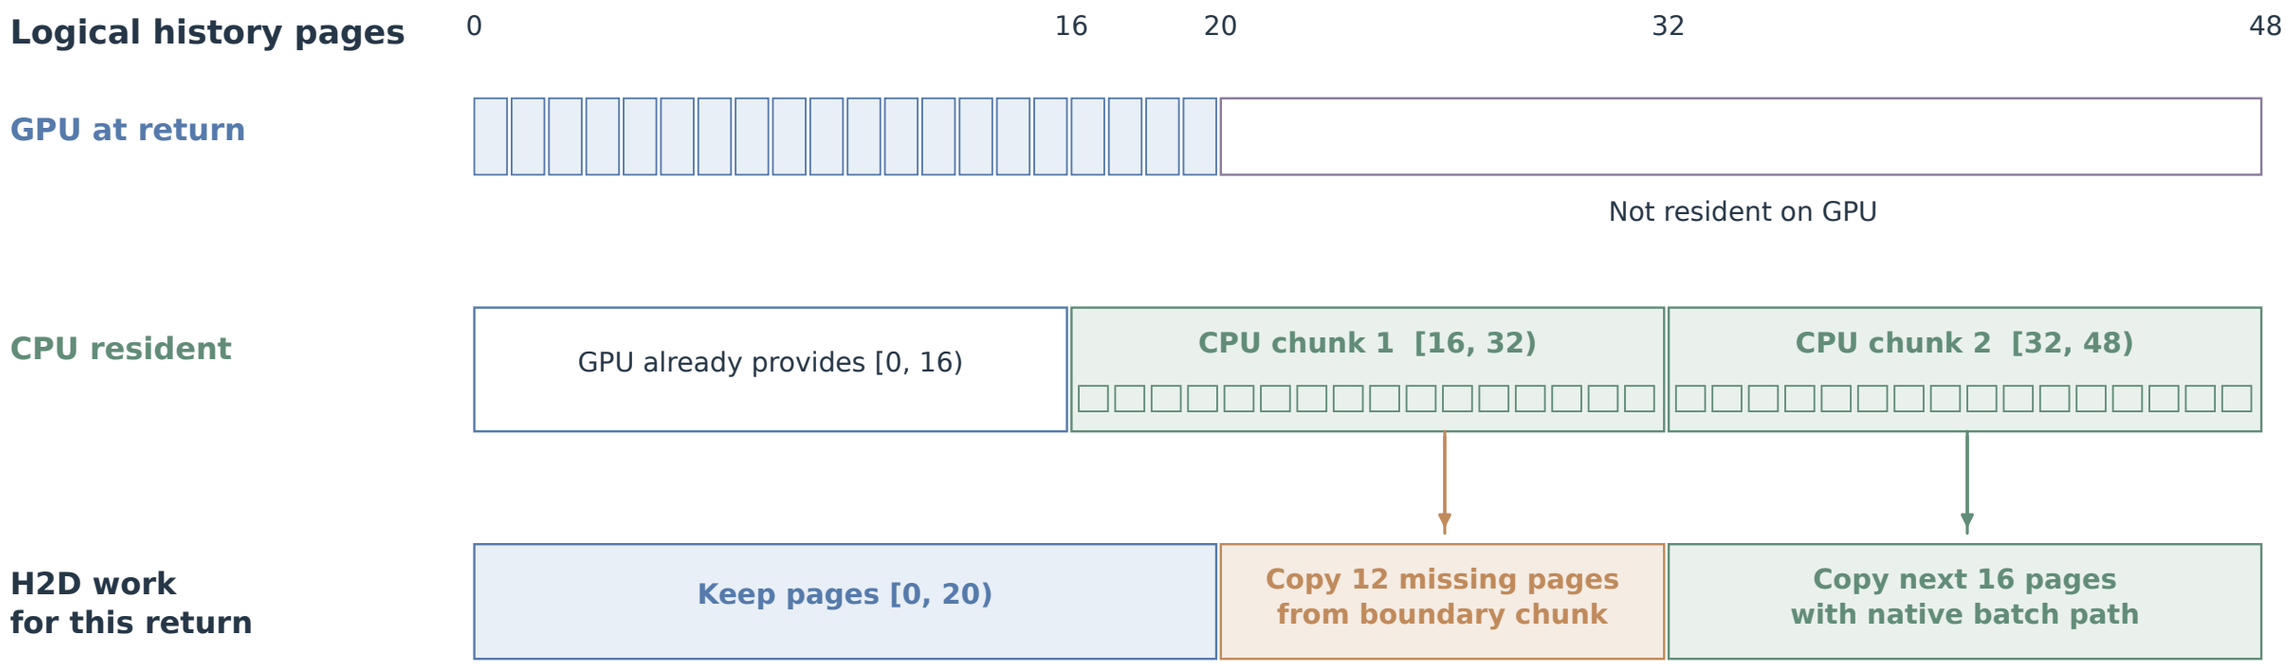

Joint recoverable endpoint: page 48; H2D transfers 28 pages, with no copy over [16, 20).

1 Join coverage

Extend the GPU prefix through consecutive, confirmed CPU chunks.

2 Respect page boundaries

Preserve resident GPU pages; map only the missing chunk suffix.

3 Retire transfer references

DMA completion precedes page reuse; recompute beyond any coverage gap.

The page strip shows logical history order, not physically contiguous GPU allocation. No transfer overlap or speedup is assumed.